

tobacco_consumption

April 10, 2022

1 Tobacco Consumption

Tobacco consumption is one of the primary causes of lung cancer in the World. Tobacco in the form of cigars and cigarettes is usually available to adult population in many supermarkets and grocery stores. The data obtained for this analysis describes Tobacco Consumption in USA from 2000 to 2020. From behavior of the data in those 21 years, the aim of the project is to predict total tobacco consumption in 2021 and 2022.

At first, the libraries used for this project are imported.

```
[ ]: import pandas as pd
import matplotlib.pyplot as plt
import statsmodels.api as sm
import numpy as np
import seaborn as sns
import random
import math
from statsmodels.tsa.seasonal import seasonal_decompose
from statsmodels.tsa.api import Holt
from statsmodels.tsa.arima.model import ARIMA
from sklearn.linear_model import LinearRegression
from sklearn.metrics import mean_squared_error, mean_absolute_percentage_error
```

An additional import is included in order to ignore some warnings while processing the data.

```
[ ]: import warnings
warnings.simplefilter(action="ignore", category=FutureWarning)
```

1.1 Extraction

The data for this project is stored in a *.csv* file. The path to the file is defined in the variable *DATA_PATH*.

```
[ ]: DATA_PATH = "../data/Tobacco_Consumption.csv"
```

The file is read and a sample of the data is shown.

```
[ ]: tobacco_data_raw = pd.read_csv(DATA_PATH)
tobacco_data_raw.sample(10)
```

[]:	Year	Location	Abbrev	Location	Desc	Population	Topic	\
151	2011		US	National		237657645	Noncombustible Tobacco	
69	2005		US	National		222003984	Noncombustible Tobacco	
207	2015		US	National		247773709	Combustible Tobacco	
134	2010		US	National		235153929	Combustible Tobacco	
148	2011		US	National		237657645	Combustible Tobacco	
177	2013		US	National		242542967	Combustible Tobacco	
209	2016		US	National		249485228	Combustible Tobacco	
59	2004		US	National		219552929	Combustible Tobacco	
229	2017		US	National		252063800	Combustible Tobacco	
96	2007		US	National		227239768	Combustible Tobacco	

	Measure	Submeasure	Data Value	Unit	\
151	Smokeless Tobacco	Snuff		Pounds	
69	Smokeless Tobacco	Chewing Tobacco		Pounds	
207	Cigars	Small Cigars		Cigars	
134	Loose Tobacco	Pipe Tobacco		Pounds	
148	Loose Tobacco	Roll-Your-Own Tobacco	Cigarette	Equivalents	
177	Loose Tobacco	Total Loose Tobacco		Pounds	
209	Cigars	Large Cigars		Cigars	
59	Loose Tobacco	Pipe Tobacco		Pounds	
229	Loose Tobacco	Roll-Your-Own Tobacco	Cigarette	Equivalents	
96	Loose Tobacco	Roll-Your-Own Tobacco	Cigarette	Equivalents	

	Domestic	Imports	Total	Domestic Per Capita	\
151	103097456	464519	103561975	0.434	
69	38883276	315916	39199192	0.175	
207	530680751	23505000	554185751	2.000	
134	22266231	2822447	25088678	0.000	
148	2412067938	209522215	2621590154	10.000	
177	42645920	3456679	46102599	0.000	
209	5056761893	6974839000	12031600893	20.000	
59	3904810	794487	4699297	0.000	
229	1212400738	99499323	1311900062	5.000	
96	8508941785	816734031	9325675815	37.000	

	Imports Per Capita	Total Per Capita
151	0.002	0.436
69	0.001	0.177
207	0.000	2.000
134	0.000	0.000
148	1.000	11.000
177	0.000	0.000
209	28.000	48.000
59	0.000	0.000
229	0.000	5.000
96	4.000	41.000

1.2 Exploratory Data Analysis

Describe data table

```
[ ]: tobacco_data_raw.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 273 entries, 0 to 272
Data columns (total 14 columns):
 #   Column                Non-Null Count  Dtype
---  -
 0   Year                  273 non-null   int64
 1   LocationAbbrev        273 non-null   object
 2   LocationDesc          273 non-null   object
 3   Population            273 non-null   int64
 4   Topic                 273 non-null   object
 5   Measure               273 non-null   object
 6   Submeasure           273 non-null   object
 7   Data Value Unit       273 non-null   object
 8   Domestic              273 non-null   int64
 9   Imports               273 non-null   int64
10   Total                 273 non-null   int64
11   Domestic Per Capita   273 non-null   float64
12   Imports Per Capita    273 non-null   float64
13   Total Per Capita      273 non-null   float64
dtypes: float64(3), int64(5), object(6)
memory usage: 30.0+ KB
```

In this table, there are categorical and numerical variables.

The exploration will initially focus on categorical variables and later on the numerical ones.

1.2.1 Categorical Data Exploration

The categorical data columns are filtered from the original dataframe.

```
[ ]: # Filter categorical variables from data
tobacco_categorical_data = tobacco_data_raw.select_dtypes(exclude=['int',
↪ 'float'])
# Show head of tables
tobacco_categorical_data.head(10)
```

```
[ ]: LocationAbbrev LocationDesc Topic Measure \
0      US      National Noncombustible Tobacco Smokeless Tobacco
1      US      National Combustible Tobacco Cigarettes
2      US      National Combustible Tobacco Cigars
3      US      National Combustible Tobacco Loose Tobacco
4      US      National Combustible Tobacco Loose Tobacco
5      US      National Combustible Tobacco Cigars
6      US      National Combustible Tobacco Loose Tobacco
```

7	US	National	Combustible Tobacco	Loose Tobacco
8	US	National	Combustible Tobacco	Loose Tobacco
9	US	National	Combustible Tobacco	Cigars

	Submeasure	Data Value Unit
0	Chewing Tobacco	Pounds
1	Cigarette Removals	Cigarettes
2	Total Cigars	Cigars
3	Total Loose Tobacco	Cigarette Equivalents
4	Total Loose Tobacco	Pounds
5	Small Cigars	Cigars
6	Pipe Tobacco	Pounds
7	Roll-Your-Own Tobacco	Cigarette Equivalents
8	Roll-Your-Own Tobacco	Pounds
9	Large Cigars	Cigars

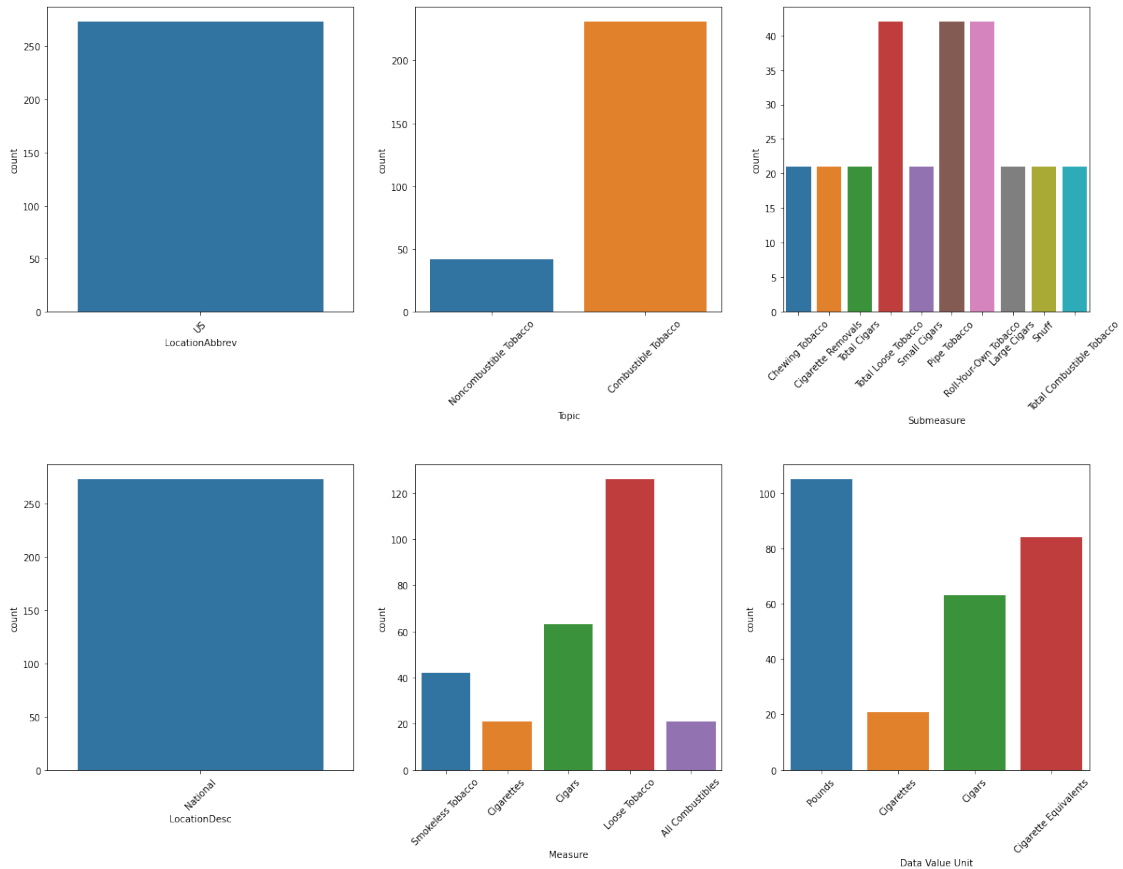
Categorical data columns are identified.

```
[ ]: # Show numbers of columns
print(f"There is a total of {len(tobacco_categorical_data.columns)}
      ↪categorical data columns")
# Show name of the columns
print(f"The columns are: {tobacco_categorical_data.columns}")
```

```
There is a total of 6 categorical data columns
The columns are: Index(['LocationAbbrev', 'LocationDesc', 'Topic', 'Measure',
                        'Submeasure',
                        'Data Value Unit'],
                        dtype='object')
```

To explore the frequency of elements for each column, frequency is plotted in a bar chart, where x axis is the name of the elements in the column, and yaxis is the number of times the element is in the column.

```
[ ]: # Create plot object
fig, ax = plt.subplots(2,3, figsize=(20, 15))
fig.subplots_adjust(hspace=.5)
i = 0
# Add subplot of frequency of elements per column of categoricall data
for col in tobacco_categorical_data.columns:
    sns.countplot(tobacco_categorical_data[col], ax=ax[i%2, math.floor(i/2)])
    i+=1
# Rotate axis of each subplot
for ax in fig.axes:
    plt.sca(ax)
    plt.xticks(rotation=45)
```



For *LocationDesc* and *LocationAbbrev* columns there is only one unique value each. Therefore, these columns are constants.

Most values in submeasure have a 21 apperances in the table.

The combinations of values in the columns “Measure”, “Submeasure” and “Units” is further explored, to identify how many time each different combinations is shown in the table.

Categorical data combinations Unique combinations of categories are obtained.

```
[ ]: # Get unique combinations by dropping duplicated categorical columns
tobacco_categorical_data.drop_duplicates()
```

```
[ ]: LocationAbbrev LocationDesc Topic Measure \
0 US National Noncombustible Tobacco Smokeless Tobacco
1 US National Combustible Tobacco Cigarettes
2 US National Combustible Tobacco Cigars
3 US National Combustible Tobacco Loose Tobacco
4 US National Combustible Tobacco Loose Tobacco
5 US National Combustible Tobacco Cigars
6 US National Combustible Tobacco Loose Tobacco
```

7	US	National	Combustible Tobacco	Loose Tobacco
8	US	National	Combustible Tobacco	Loose Tobacco
9	US	National	Combustible Tobacco	Cigars
10	US	National	Combustible Tobacco	Loose Tobacco
11	US	National	Noncombustible Tobacco	Smokeless Tobacco
12	US	National	Combustible Tobacco	All Combustibles

	Submeasure	Data Value	Unit
0	Chewing Tobacco		Pounds
1	Cigarette Removals		Cigarettes
2	Total Cigars		Cigars
3	Total Loose Tobacco	Cigarette	Equivalents
4	Total Loose Tobacco		Pounds
5	Small Cigars		Cigars
6	Pipe Tobacco		Pounds
7	Roll-Your-Own Tobacco	Cigarette	Equivalents
8	Roll-Your-Own Tobacco		Pounds
9	Large Cigars		Cigars
10	Pipe Tobacco	Cigarette	Equivalents
11	Snuff		Pounds
12	Total Combustible Tobacco	Cigarette	Equivalents

Describe combinations and unique combinations.

```
[ ]: # Get number of unique combinations and total combinations in the table
total_categories_combinations = len(tobacco_categorical_data)
unique_categories_combinations = len(tobacco_categorical_data.drop_duplicates())
# Print summary
print(f"Total combinations of categories (rows):␣
↪{total_categories_combinations}")
print(f"Find {unique_categories_combinations} unique category combinations")
print(f"Relation: {total_categories_combinations/
↪unique_categories_combinations}")
```

Total combinations of categories (rows): 273

Find 13 unique category combinations

Relation: 21.0

13 combinations are repeated 21 times in the table.

This number match the number of years in the data. The dataset included 13 different values per year.

1.2.2 Numerical Data Exploration

The numerical data columns are filtered from the original dataframe.

```
[ ]:
```

```
# Filter numerical variables from data
tobacco_numerical_data = tobacco_data_raw.select_dtypes(include=['int',
↳ 'float'])
# Show head of tables
tobacco_numerical_data.head(10)
```

```
[ ]:      Year  Population      Domestic      Imports      Total \
0  2000   209786736   45502156      91965   45594121
1  2000   209786736  423250355675  12319663000  435570018675
2  2000   209786736   5612867329   548243000   6161110329
3  2000   209786736   8291276800   702741662   8994018462
4  2000   209786736   16841656    1427444    18269100
5  2000   209786736   2243135044   36049000   2279184044
6  2000   209786736    5352683    739887    6092570
7  2000   209786736   5656109785   338489600   5994599385
8  2000   209786736    11488973    687557    12176530
9  2000   209786736   3369732285   512194000   3881926285

      Domestic Per Capita  Imports Per Capita  Total Per Capita
0                0.217                0.0                0.217
1            2018.000                59.0            2076.000
2              27.000                3.0              29.000
3              40.000                3.0              43.000
4               0.000                0.0               0.000
5              11.000                0.0              11.000
6               0.000                0.0               0.000
7              27.000                2.0              29.000
8               0.000                0.0               0.000
9              16.000                2.0              19.000
```

Numerical data columns are identified.

```
[ ]: # Show numbers of columns
print(f"There is a total of {len(tobacco_numerical_data.columns)} numerical_
↳ data columns")
# Show name of the columns
print(f"The columns are: {tobacco_numerical_data.columns}")
```


There is a total of 8 numerical data columns

The columns are: Index(['Year', 'Population', 'Domestic', 'Imports', 'Total',
'Domestic Per Capita', 'Imports Per Capita', 'Total Per Capita'],
dtype='object')

To understand how each variable is related to each other, correlations are obtained and plotted.

```
[ ]: # Explore correlations
correlations = tobacco_numerical_data.corr()
# Plot correlations
```

```
sns.heatmap(correlations, annot=True)
plt.show()
```



Year and *Population* have a strong correlation with each other, but a low correlation to tobacco values.

Per capita values have a strong correlation with normal values.

A test is applied to verify if per capita values are obtained from total values and population.

```
[ ]: # Obtain difference between per capital columns and normal columns divided by
      ↪ population
relation_per_capita = round(tobacco_numerical_data["Total"] /
      ↪ tobacco_numerical_data["Population"], 1) - tobacco_numerical_data["Total Per_
      ↪ Capita"]
round(relation_per_capita.median(), 3)
```

```
[ ]: 0.003
```

The difference is close to 0. Therefore, the next expressions can be established from the data:

$$\text{Domestic_per_capita} = \frac{\text{Domestic}}{\text{Population}}$$

$$\text{Imports_per_capita} = \frac{\text{Imports}}{\text{Population}}$$

$$\text{Total_per_capita} = \frac{\text{Total}}{\text{Population}}$$

For further analysis, per capita columns are excluded.

Domestic and *Imports* have a strong correlation to *Total* column.

```
[ ]: # Difference between total and imports + domestic is obtained
difference_total = tobacco_numerical_data["Total"] -
↳ tobacco_numerical_data["Domestic"] - tobacco_numerical_data["Imports"]
difference_total.median()
```

```
[ ]: 0.0
```

The difference is 0, so

$$\text{Total} = \text{Imports} + \text{Domestic}$$

To have a better understanding of these variables, it is needed to combine numerical exploration with the unique categories exploration. After that analysis, the relation between submeasures is expected to be identified.

1.2.3 Integrated Exploration (Categories & Numerical Data)

As each year has the same category combinations, one year (2000) is used as a sample. As this analysis focuses in tobacco consumption, only units related to products are taken in to account. Therefore, unit “Pounds” is excluded.

```
[ ]: # Get filtered df
products_df = tobacco_data_raw[(tobacco_data_raw["Data Value Unit"] !=
↳ "Pounds") & (tobacco_data_raw["Year"] == 2000)]
products_df
```

```
[ ]:
   Year LocationAbbrev LocationDesc Population Topic \
1   2000      US      National   209786736 Combustible Tobacco
2   2000      US      National   209786736 Combustible Tobacco
3   2000      US      National   209786736 Combustible Tobacco
5   2000      US      National   209786736 Combustible Tobacco
7   2000      US      National   209786736 Combustible Tobacco
9   2000      US      National   209786736 Combustible Tobacco
10  2000      US      National   209786736 Combustible Tobacco
12  2000      US      National   209786736 Combustible Tobacco

   Measure              Submeasure      Data Value Unit \
1   Cigarettes      Cigarette Removals      Cigarettes
2   Cigars          Total Cigars          Cigars
3   Loose Tobacco      Total Loose Tobacco  Cigarette Equivalents
5   Cigars          Small Cigars          Cigars
7   Loose Tobacco      Roll-Your-Own Tobacco  Cigarette Equivalents
```

	Cigars		Large Cigars		Cigars
	Loose Tobacco		Pipe Tobacco		Cigarette Equivalents
12	All Combustibles	Total Combustible Tobacco			Cigarette Equivalents

	Domestic	Imports	Total	Domestic Per Capita \
1	423250355675	12319663000	435570018675	2018.0
2	5612867329	548243000	6161110329	27.0
3	8291276800	702741662	8994018462	40.0
5	2243135044	36049000	2279184044	11.0
7	5656109785	338489600	5994599385	27.0
9	3369732285	512194000	3881926285	16.0
10	2635167015	364252062	2999419077	13.0
12	437154499804	13570647662	450725147466	2084.0

	Imports Per Capita	Total Per Capita
1	59.0	2076.0
2	3.0	29.0
3	3.0	43.0
5	0.0	11.0
7	2.0	29.0
9	2.0	19.0
10	2.0	14.0
12	65.0	2148.0

Total Loose Tobacco is compared to Pipe Tobacco and Roll-Your-Own Tobacco

```
[ ]: # Compare diff between Total Loose Tobacco and Pipe Tobacco
products_df["Domestic"][3] - products_df["Domestic"][7] -
↳products_df["Domestic"][10]
```

[]: 0

Therefore,

$$Total_Loose_Tobacco = Pipe_Tobacco + Roll_Your_Own_Tobacco$$

The Loose Tobacco values are in the table twice (as pounds and as cigarette equivalents), that's the reason the frequency was the double than other cases in categorical data analysis.

Total Cigars are compared to Small and Large Cigars...

```
[ ]: products_df["Domestic"][2] - products_df["Domestic"][5] -
↳products_df["Domestic"][9]
```

[]: 0

For Cigars:

$$Total_Cigars = Small_Cigars + Large_Cigars$$

```
[ ]: # Get sum of non-total submeasures
sum_cigarettes = products_df["Domestic"][~products_df["Submeasure"].str.
    ↪contains("Total")].sum()
# Compare sum to Total Combustible Tobacco variable
products_df["Domestic"][products_df["Submeasure"]=="Total Combustible Tobacco"]_
    ↪- sum_cigarettes
```

```
[ ]: 12    0
      Name: Domestic, dtype: int64
```

$Total_Combustible_Tobacco = Total_Cigars + Total_Loose_Tobacco + Cigarette_Removals$

Cigarette, Cigarette Equivalents, and Cigars units have a 1:1:1 relationship.

Total Combustible Tobacco contains information of all types of tobacco products submeasures. This value will be the target variable that is going to be predicted in the analysis.

1.3 Data Wrangling

The original dataframe is filtered and transformed to get a useful table focused in the target variable. Unnecessary columns are drop and year is set as index of the table.

```
[ ]: total_combustible_tobacco_df =_
    ↪tobacco_data_raw[tobacco_data_raw["Submeasure"]=="Total Combustible Tobacco"]
# Drop columns with constant information
total_combustible_tobacco_df.drop(columns=["LocationAbbrev", "LocationDesc",_
    ↪"Topic", "Measure",
    ↪"Submeasure", "Data Value Unit"], inplace=True)
# To reduce data with similar behavior, per capita values will be also ignored_
    ↪in the transformation
total_combustible_tobacco_df.drop(columns=["Domestic Per Capita", "Imports Per_
    ↪Capita", "Total Per Capita"], inplace=True)
# Year to index and datetime object
total_combustible_tobacco_df.set_index("Year", inplace = True)
total_combustible_tobacco_df.index = pd.
    ↪to_datetime(total_combustible_tobacco_df.index, format = "%Y")
# Show time series
total_combustible_tobacco_df
```

C:\Users\edson\AppData\Local\Temp\ipykernel_10308\332722847.py:3:
 SettingWithCopyWarning:
 A value is trying to be set on a copy of a slice from a DataFrame

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy

```
total_combustible_tobacco_df.drop(columns=["LocationAbbrev", "LocationDesc",
"Topic", "Measure",
```

```
C:\Users\edson\AppData\Local\Temp\ipykernel_10308\332722847.py:6:
SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame
```


See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy

```
total_combustible_tobacco_df.drop(columns=["Domestic Per Capita", "Imports Per
Capita", "Total Per Capita"], inplace=True)
```

```
[ ]:      Population      Domestic      Imports      Total
Year
2000-01-01  209786736  437154499804  13570647662  450725147466
2001-01-01  212345162  424447310483  16245487123  440692797606
2002-01-01  214754648  408303465673  22459975923  430763441596
2003-01-01  217068101  391196952717  24733051892  415930004609
2004-01-01  219552929  390166439041  24254733538  414421172579
2005-01-01  222003984  380712422031  20474781600  401187203631
2006-01-01  224583123  382748665603  18493161400  401241827003
2007-01-01  227239768  367880476823  16206169554  384086646377
2008-01-01  229945137  357080696204  14183410908  371264347112
2009-01-01  232458335  329491123539  12632961400  342124084939
2010-01-01  235153929  317475052695  11764157615  329239210311
2011-01-01  237657645  315683514293  10893776123  326577290416
2012-01-01  240185952  310009439421  12386990662  322396430083
2013-01-01  242542967  295476582903  14164450662  309641033564
2014-01-01  245273438  282020020624  16175868538  298195889163
2015-01-01  247773709  285537129046  17214155385  302751284431
2016-01-01  249485228  273881993213  17154801323  291036794536
2017-01-01  252063800  263359518483  17581121431  280940639914
2018-01-01  253768092  249351495626  19823602569  269167785869
2019-01-01  255200373  234601291857  21813751754  256415043611
2020-01-01  256662010  235888722742  21726898662  257615621404
```

Plot total over the years

```
[ ]: sns.lineplot(x=total_combustible_tobacco_df.index,
    y=total_combustible_tobacco_df["Total"])
plt.title("Total Combustible Tobacco per Year")
plt.ylabel("Units")
plt.show()
```



Store new table as csv.

```
[ ]: # Export ts to df
OUTPUT_PATH = "../data/Transformed_Tobacco_Consumption.csv"
total_combustible_tobacco_df.to_csv(OUTPUT_PATH, index=False)
```

1.4 Exploration of Transformed Data

```
[ ]: # Explore variables distribution
total_combustible_tobacco_df.describe().convert_dtypes()
```

```
[ ]:
```

	Population	Domestic	Imports \
count	21.0	21.0	21.0
mean	234547860.285714	330117467277.190491	17331140748.761906
std	15123590.490099	63129184488.71344	4140066674.594648
min	209786736.0	234601291857.0	10893776123.0
25%	222003984.0	282020020624.0	14164450662.0
50%	235153929.0	317475052695.0	17154801323.0
75%	247773709.0	382748665603.0	20474781600.0
max	256662010.0	437154499804.0	24733051892.0

	Total
count	21.0
mean	347448271248.571411
std	63456602238.587479

```

min          256415043611.0
25%          298195889163.0
50%          329239210311.0
75%          401241827003.0
max          450725147466.0

```


Show boxplots and histograms

```

[ ]: # Create boxplots
COLORS = ["b", "g", "r", "c", "m", "y"]
fig, ax = plt.subplots(2, 2, figsize=(10, 7))
i = 0
for col in total_combustible_tobacco_df.columns:
    sns.boxplot(y=col, data=total_combustible_tobacco_df, color = random.
    choice(COLORS), ax=ax[i%2, math.floor(i/2)])
    i+=1

plt.tight_layout()

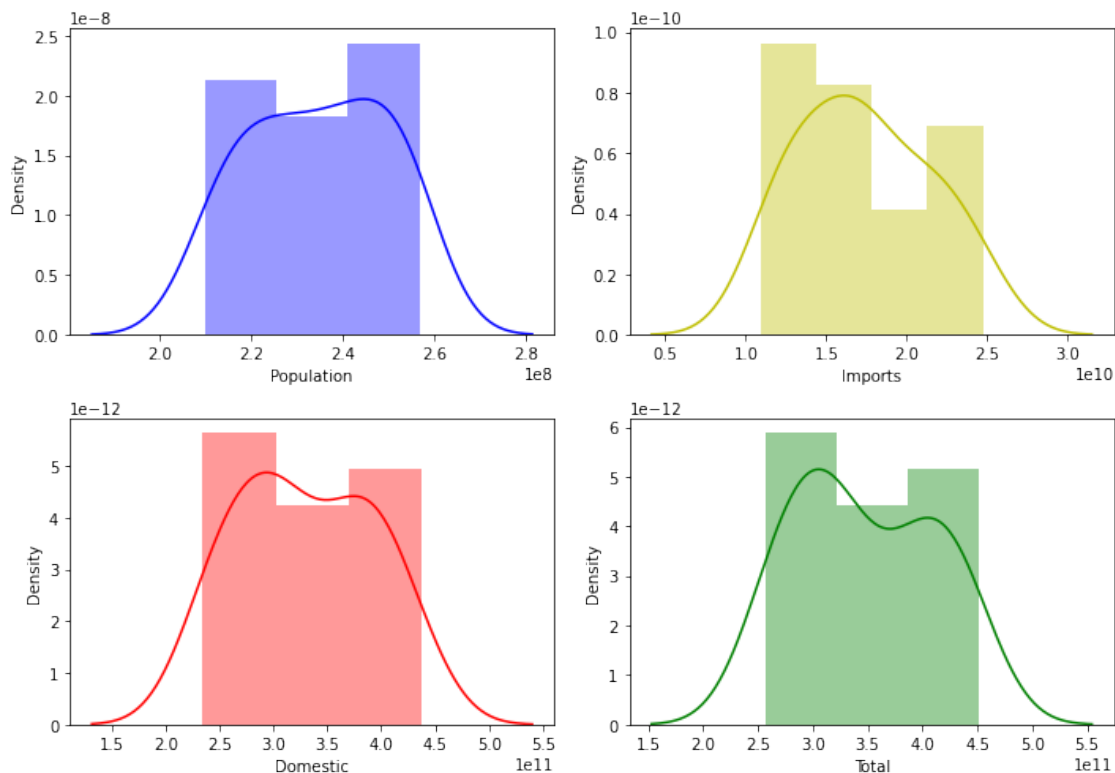
```



For population and total, data seems to be symmetric. However, Domestic and Total are a little right-skewed.

```
[ ]: # Create histogram
COLORS = ["b", "g", "r", "c", "m", "y"]
fig, ax = plt.subplots(2, 2, figsize=(10, 7))
i = 0
for col in total_combustible_tobacco_df.columns:
    sns.distplot(total_combustible_tobacco_df[col], color = random.
    choice(COLORS), ax=ax[i%2, math.floor(i/2)])
    i+=1

plt.tight_layout()
```



All variables seem close to be symmetric. The previously identified as skewed variables are also close to the center of the data.

Plots all trends by year.

```
[ ]: COLORS = ["b", "g", "r", "c", "m", "y"]
fig, ax = plt.subplots(2, 2, figsize=(10, 7))
i = 0
for col in total_combustible_tobacco_df.columns:
    sns.lineplot(x=total_combustible_tobacco_df.index,
    y=total_combustible_tobacco_df[col],
    color = random.choice(COLORS), ax=ax[i%2, math.floor(i/2)])
```

```
i+=1
plt.tight_layout()
```



The percentage of change of variables over time is explored.

1.4.1 % Change over the years

Get % of change of each variable and plot.

```
[ ]: ts_change_df = total_combustible_tobacco_df.pct_change().dropna()
ts_change_df = round(ts_change_df *100,2)
ts_change_df.head(5)
```

```
[ ]:
      Population  Domestic  Imports  Total
Year
2001-01-01      1.22     -2.91    19.71  -2.23
2002-01-01      1.13     -3.80    38.25  -2.25
2003-01-01      1.08     -4.19    10.12  -3.44
2004-01-01      1.14     -0.26     -1.93  -0.36
2005-01-01      1.12     -2.42   -15.58  -3.19
```


```
[ ]: # Plot % of change of variables
COLORS = ["b", "g", "r", "c", "m", "y"]
```



```

fig, ax = plt.subplots(2,2, figsize=(10,7))
i = 0
for col in ts_change_df.columns:
    sns.lineplot(x=ts_change_df.index, y=ts_change_df[col], color = random.
    choice(COLORS), ax=ax[i%2, math.floor(i/2)])
    i+=1
plt.tight_layout()

```



There is no clear behavior related to how much does each variables changes per year.

1.4.2 Stationarity

Augmented Dickey-Fuller test is applied to verify if the data is stationary.

```

[ ]: X = total_combustible_tobacco_df["Total"].values

result = sm.tsa.stattools.adfuller(X)
print('ADF Statistic: %f' % result[0])
print('p-value: %f' % result[1])
print('Critical Values:')
for key, value in result[4].items():
    print('\t%s: %.3f' % (key, value))


```

```
ADF Statistic: -4.365042
p-value: 0.000342
Critical Values:
    1%: -4.138
    5%: -3.155
   10%: -2.714
```


The p-value is really small (less than 5% threshold), so it is confirmed the total column is a stationary time series.

1.4.3 ACF and PACF

```
[ ]: plt.rc("figure", figsize=(10,6))
sm.graphics.tsa.plot_acf(total_combustible_tobacco_df["Total"], lags=20)
plt.show()
```




```
[ ]: plt.rc("figure", figsize=(10,6))
sm.graphics.tsa.plot_pacf(total_combustible_tobacco_df["Total"], lags=9)
plt.show()
```



1.4.4 Time Series Decomposition

```
[ ]: ts_decompose = seasonal_decompose(total_combustible_tobacco_df["Total"],  
    ↪ model="additive")  
ts_decompose.plot()  
plt.show()
```



There is neither seasonal component or resid.

1.5 Modeling

In this project, three models are compared: - Linear Regression - AutoRegresive Integrated Moving Average (ARIMA) - Holt's Exponential Smoothing

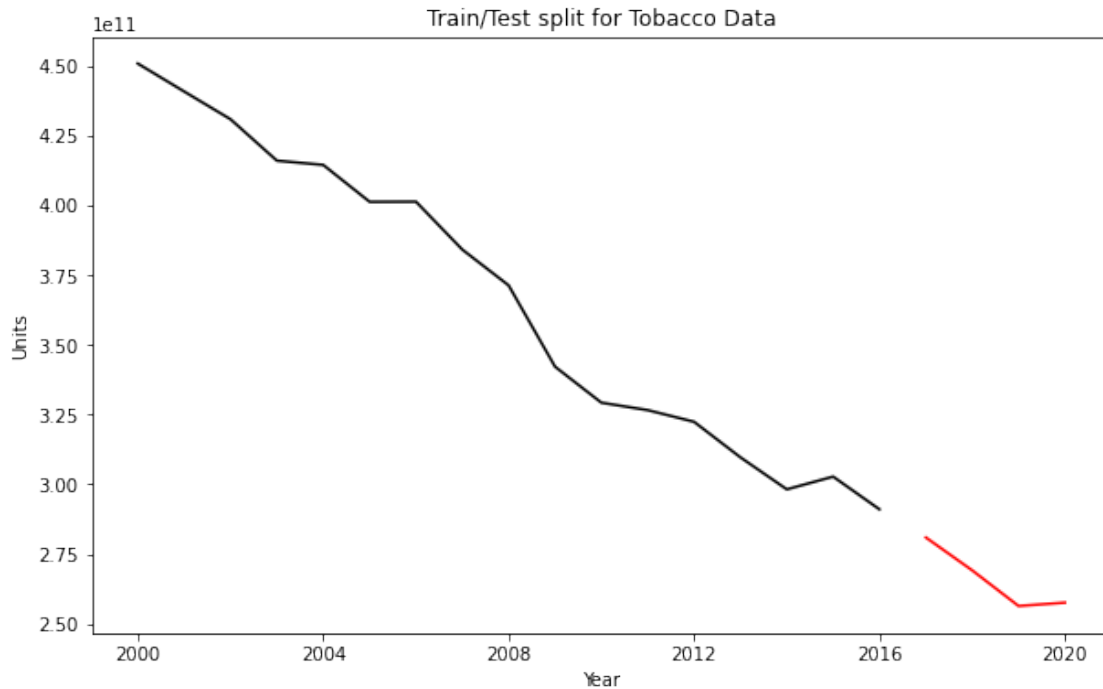
Two metrics are used for comparing and selecting a model: - Relative Root Mean Square Error (rRMSE) - Mean Absolute Percentage Error (MAPE)

One of the models will be selected to predict the total tobacco consumption over the next years.

At first, data is splitted in test and train datasets

```
[ ]: train = total_combustible_tobacco_df["Total"][total_combustible_tobacco_df.
        ↪index < "12-12-2016"]
test = total_combustible_tobacco_df["Total"][total_combustible_tobacco_df.index_
        ↪ > "12-12-2016"]

plt.plot(train, color="black")
plt.plot(test, color = "red" )
plt.ylabel("Units")
plt.xlabel("Year")
plt.title("Train/Test split for Tobacco Data")
plt.show()
```



1.5.1 Linear Regression Model

In this model, independent variable is Year and dependent variable is the total tobacco consumption. Based on that, X and Y are defined.

```
[ ]: # Get X and Y from training data
X = train.index.year.values.reshape(-1, 1)
Y = train.values.reshape(-1, 1)
```

The model is created and trained.

```
[ ]: LR_model = LinearRegression()
LR_model.fit(X, Y)
```

```
[ ]: LinearRegression()
```


To test the model, predictions of the years in the test set (2017-2020) are made.

```
[ ]: y_pred_lr = pd.Series(LR_model.predict(test.index.year.values.reshape(-1, 1)).
    ↪flatten(), index=test.index)
y_pred_lr
```

```
[ ]: Year
2017-01-01    2.706961e+11
2018-01-01    2.600396e+11
2019-01-01    2.493831e+11
```

```
2020-01-01    2.387266e+11
dtype: float64
```

```
[ ]: plt.plot(train, color="black", label = "Train")
plt.plot(test, color = "red", label = "Test")
plt.plot(y_pred_lr, color = "blue", label= "Linear Regression")
plt.ylabel("Units consumed")
plt.xlabel("Year")
plt.title("Linear Regression prediction for Tobacco Data")
plt.legend(loc="upper right")
plt.show()
```



1.5.2 ARIMA Model

Train model and tune hyperparameters.

```
[ ]: ARIMA_model = ARIMA(train, order=(3,3,2))
ARIMA_model = ARIMA_model.fit()
print(ARIMA_model.summary())
```

```
c:\Users\edson\Documents\Repositorios\Data-Science\Tobacco-Consumption-
Prediction\code\jenv\lib\site-packages\statsmodels\tsa\base\tsa_model.py:471:
ValueWarning: No frequency information was provided, so inferred frequency AS-
JAN will be used.
```

```
self._init_dates(dates, freq)
```

```

c:\Users\edson\Documents\Repositorios\Data-Science\Tobacco-Consumption-
Prediction\code\jenv\lib\site-packages\statsmodels\tsa\base\tsa_model.py:471:
ValueWarning: No frequency information was provided, so inferred frequency AS-
JAN will be used.
    self._init_dates(dates, freq)
c:\Users\edson\Documents\Repositorios\Data-Science\Tobacco-Consumption-
Prediction\code\jenv\lib\site-packages\statsmodels\tsa\base\tsa_model.py:471:
ValueWarning: No frequency information was provided, so inferred frequency AS-
JAN will be used.
    self._init_dates(dates, freq)
c:\Users\edson\Documents\Repositorios\Data-Science\Tobacco-Consumption-
Prediction\code\jenv\lib\site-
packages\statsmodels\tsa\statespace\sarimax.py:978: UserWarning: Non-invertible
starting MA parameters found. Using zeros as starting parameters.
    warn('Non-invertible starting MA parameters found.')

```

SARIMAX Results

```

=====
Dep. Variable:                Total    No. Observations:                17
Model:                      ARIMA(3, 3, 2)    Log Likelihood                -341.950
Date:                      Sun, 10 Apr 2022    AIC                695.901
Time:                      12:34:09    BIC                699.735
Sample:                      01-01-2000    HQIC                695.546
                        - 01-01-2016

```

Covariance Type: opg

```

=====
              coef      std err          z      P>|z|      [0.025      0.975]
-----
ar.L1         -1.0102      0.479      -2.108      0.035      -1.950      -0.071
ar.L2         -0.2133      0.257      -0.830      0.407      -0.717      0.291
ar.L3         -0.0263      0.088      -0.301      0.763      -0.198      0.145
ma.L1         -0.2394      0.612      -0.391      0.696      -1.439      0.960
ma.L2         -0.7161      0.851      -0.842      0.400      -2.383      0.951
sigma2        1.172e+20    7.68e-21    1.53e+40    0.000      1.17e+20    1.17e+20
=====

```

```

===
Ljung-Box (L1) (Q):                0.25    Jarque-Bera (JB):
0.83
Prob(Q):                0.62    Prob(JB):
0.66
Heteroskedasticity (H):                6.17    Skew:
0.53
Prob(H) (two-sided):                0.07    Kurtosis:
2.44
=====
===

```

Warnings:

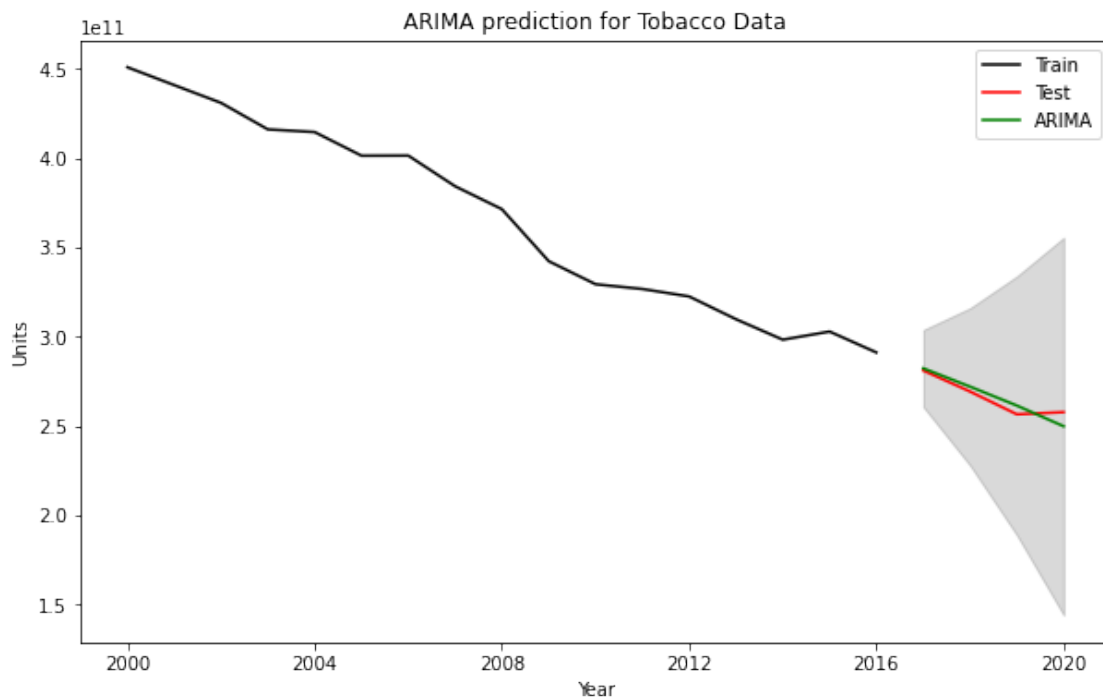
[1] Covariance matrix calculated using the outer product of gradients (complex-step).

[2] Covariance matrix is singular or near-singular, with condition number $2.06e+56$. Standard errors may be unstable.

To test the model, predictions of the years in the test set (2017-2020) are made.

```
[ ]: arima_pred = ARIMA_model.get_forecast(len(test))
      # Get confidence interval
      y_conf_int_df = arima_pred.conf_int(alpha=0.05)
      y_conf_int_df
      # Get predictions for test set years
      y_pred_arima = ARIMA_model.predict(start=test.index[0], end=test.index[-1])

[ ]: plt.plot(train, color="black", label = "Train")
      plt.plot(test, color = "red", label = "Test")
      plt.plot(y_pred_arima, color = "green", label= "ARIMA")
      plt.fill_between(y_conf_int_df.index, y_conf_int_df["lower Total"],
      ↪ y_conf_int_df["upper Total"], color="k", alpha=0.15)
      plt.ylabel("Units")
      plt.xlabel("Year")
      plt.title("ARIMA prediction for Tobacco Data")
      plt.legend(loc="upper right")
      plt.show()
```



1.5.3 Holt's Exponential Smoothing Model (ETS)

Train model

```
[ ]: holt_model = Holt(train, initialization_method="estimated").  
      ↪fit(smoothing_level=0.8, smoothing_trend=0.2, optimized=False)
```

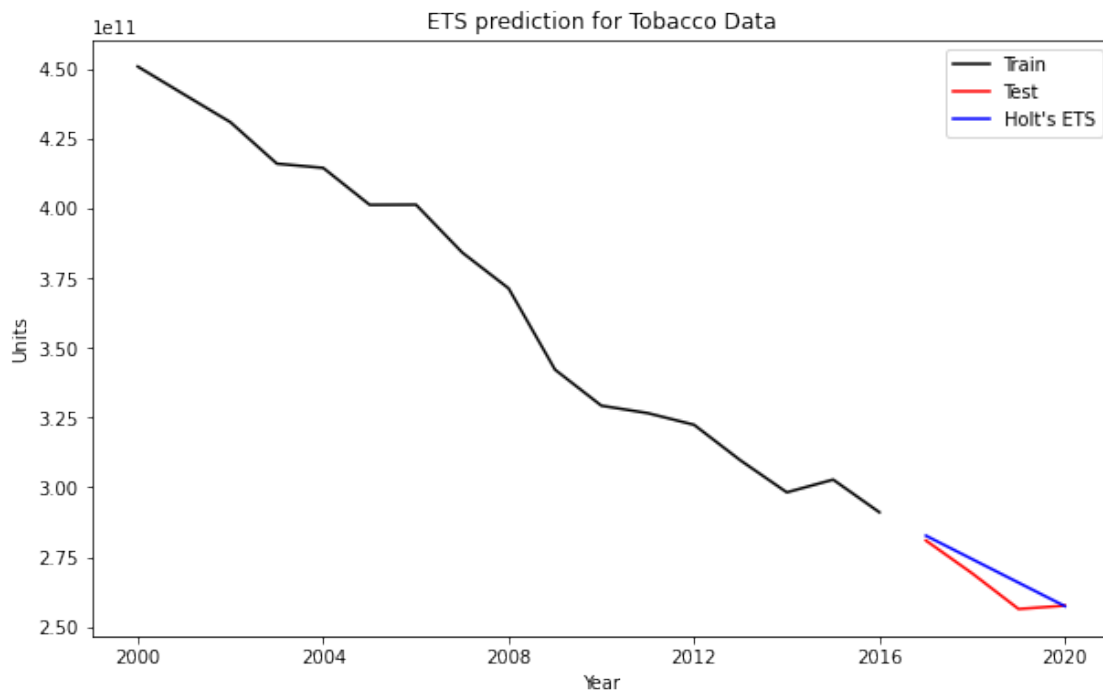
```
c:\Users\edson\Documents\Repositorios\Data-Science\Tobacco-Consumption-  
Prediction\code\jenv\lib\site-packages\statsmodels\tsa\base\tsa_model.py:471:  
ValueWarning: No frequency information was provided, so inferred frequency AS-  
JAN will be used.
```

```
self._init_dates(dates, freq)
```

To test the model, predictions of the years in the test set (2017-2020) are made.

```
[ ]: y_pred_ets = holt_model.forecast(len(test))
```

```
[ ]: plt.plot(train, color="black", label = "Train")  
      plt.plot(test, color = "red", label = "Test")  
      plt.plot(y_pred_ets, color = "blue", label= "Holt's ETS")  
      plt.ylabel("Units")  
      plt.xlabel("Year")  
      plt.title("ETS prediction for Tobacco Data")  
      plt.legend(loc="upper right")  
      plt.show()
```



1.5.4 Evaluation and Selection

Define function to obtain relative root mean square error of predicted values.

```
[ ]: def rRMSE(actual_values, predicted_values, mean_value):  
    rmse_value = np.sqrt(mean_squared_error(actual_values, predicted_values))  
    rrmse_value = rmse_value/mean_value  
    return rrmse_value
```

Create table that summarizes results of evaluation metrics.

```
[ ]: # Get evaluation metrics in the form of a dict  
data_results = {"rRSME": [rRMSE(test.values, y_pred_arma.values, test.mean()),  
                           rRMSE(test.values, y_pred_lr.values, test.mean()),  
                           rRMSE(test.values, y_pred_ets.values, test.mean())],  
               "MAPE": [mean_absolute_percentage_error(test.values, y_  
↳ y_pred_arma.values),  
                           mean_absolute_percentage_error(test.values, y_pred_lr.  
↳ values),  
                           mean_absolute_percentage_error(test.values, y_pred_ets.  
↳ values)]}  
# Dict to df  
evaluation_df = pd.DataFrame(data_results, index= ["ARIMA", "Linear_  
↳ Regression", "Holt's ETS"])  
evaluation_df
```


```
[ ]: 

|                   | rRSME    | MAPE     |
|-------------------|----------|----------|
| ARIMA             | 0.018342 | 0.016066 |
| Linear Regression | 0.045826 | 0.042781 |
| Holt's ETS        | 0.020407 | 0.015655 |


```

```
[ ]: evaluation_df.plot.bar()
```

```
[ ]: <AxesSubplot:>
```



Linear Regression model provided a worse performance than the other two models. ARIMA and ETS got similar results in MAPE metric, but performance of ARIMA was better than ETS in rRMSE metric. For ARIMA, both metric results are under 2%.

ARIMA model is selected.

1.6 Results

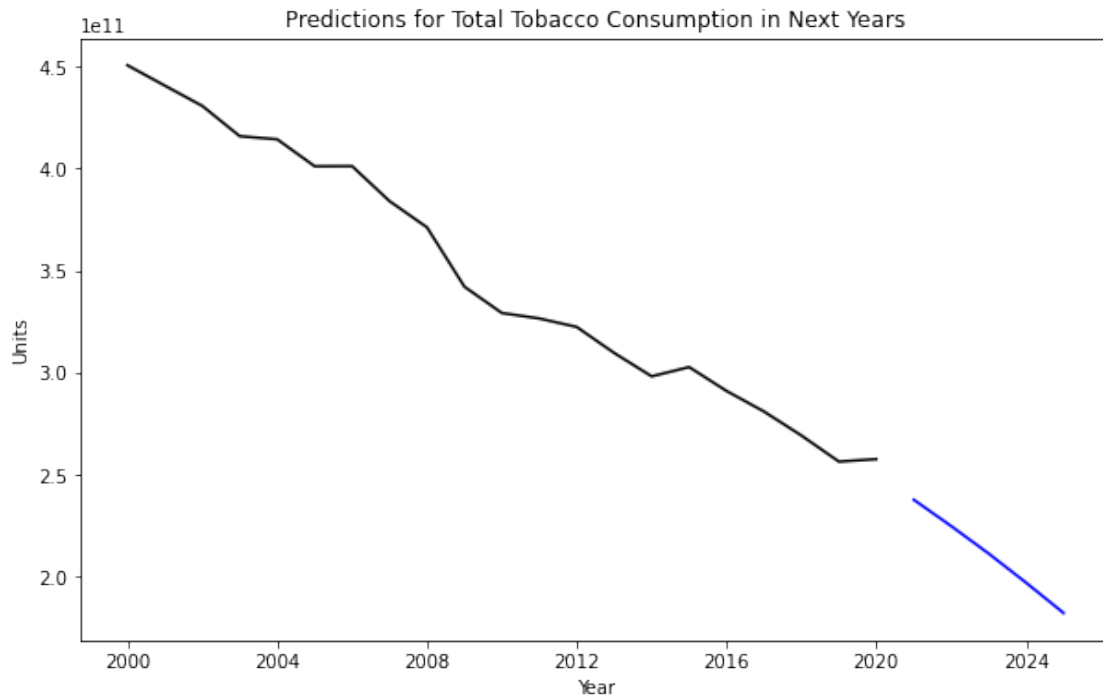
Get tobacco consumption for 2021, 2022 and 2023.

```
[ ]: y_predictions_future = ARIMA_model.predict(start="2021", end="2025")
y_predictions_future
```

```
[ ]: 2021-01-01    2.376935e+11
      2022-01-01    2.247738e+11
      2023-01-01    2.113127e+11
      2024-01-01    1.970416e+11
      2025-01-01    1.821700e+11
      Freq: AS-JAN, Name: predicted_mean, dtype: float64
```

```
[ ]: plt.plot(total_combustible_tobacco_df["Total"], color="black", label = "Train")
      plt.plot(y_predictions_future, color = "blue", label= "Predictions")
      plt.ylabel("Units")
```

```
plt.xlabel("Year")
plt.title("Predictions for Total Tobacco Consumption in Next Years")
plt.show()
```



1.7 Conclusions

Total Tobacco Consumption is decreasing through the years in the United States.

For the next years, the total consumption is expected to keep decreasing.